

ARCHITECTURE_REFERENCE

Builders Challenge Reference

A deeper dive into the architecture behind [BUILD_DAY.md](#) — the models, the DataEngine pipeline, and the full video corpus. Team setup and credentials are covered there; this doc doesn't repeat them.

Models (NVIDIA on CoreWeave)

Inference for the pipeline runs on shared **GPU endpoints** (NVIDIA Cosmos + YOLO + Canary). You don't deploy the models yourself.

Model	Role in VSS	Endpoint (env var, in /config/<team>.config)
NVIDIA Cosmos3-Reason (nvidia/cosmos3-reason)	Video understanding / reasoning over segments	\$COSMOS3_REASON_URL
YOLO11 (yolo11s , Ultralytics)	Object detection (bounding boxes / labels)	\$YOLO_URL
NVIDIA Cosmos Embed1 (nvidia/cosmos-embed1)	Text and visual embeddings (256-dim) for hybrid search	\$COSMOS_EMBED1_URL
NVIDIA Canary-1B (nvidia/canary-1b)	Speech-to-text (ASR) + speech translation — not Cosmos; NeMo audio model	\$CANARY_1B_URL

Each model has its own host and port; don't assume they share a host. No auth token is needed to call them. Ask Cursor / see `.cursor/skills/gpu/` for how to call each model.


Ingest and search call **Cosmos3-Reason**, **YOLO11**, and **Embed1** through your pipeline and backend today.

Canary-1B is not wired into the current VSS pipeline. If you want audio transcripts, spoken-word search, or speech translation in your demo — let your imagination go and hook it in. **Ask Cursor how to use Canary-1B** against your stack and skills.


Data Engine

VSS Blueprint

Your team’s ingest runs as a **VAST DataEngine** serverless pipeline. A video chunk lands in S3, then functions run in sequence until searchable rows exist in VastDB.



A separate **events / prompt-suggester** function runs on a schedule and feeds UI suggestions.



VAST DataEngine (pipeline’s tab) — view the pipeline’s flow, logs, and traces. Treat the graph as given for the challenge (re-ingest / upload via skills; don’t rebuild functions or redeploy the pipeline).

Function	What it does
Segmenter	Splits each uploaded chunk into short fixed-length clips and writes them to the segments bucket. Builders challenge: organizers already used the Segmenter

Function	What it does
	to pre-ingest your corpus. During the challenge you only re-ingest data that is already segmented and indexed — the Segmenter is not in the path you run.
Detector	Runs YOLO11 on each segment and records object classes, counts, and bbox sidecars.
Reasoner	Calls NVIDIA Cosmos3-Reason to write a searchable natural-language description of the segment.
Embedder	Calls NVIDIA Cosmos Embed1 to build text (and visual) vectors for hybrid search.
VastDB writer	Persists embeddings, reasoning, detections, and metadata as a row in your VastDB collection.
Events (prompt-suggester)	Periodically scans recent segments and writes suggested search prompts / key events for the UI.

Builders challenge note: your archive is **pre-ingested** (Segmenter already ran). Your live path is **re-ingest** on existing segments: Detector → Reasoner → Embedder → VastDB writer (via `ingest/reingest-videos / reingest-chunk`). You don't upload new chunks or invoke the Segmenter for this challenge.

You don't need to redeploy this graph for the hackathon — treat it as the engine behind search, dashboard, suggestions, and re-ingest.

What you were given

For your team you already have:

Piece	What it is
Ingest pipeline	DataEngine graph that pre-ingested your corpus (Segmenter already ran). During the challenge you re-ingest existing segments (detect → reason → embed → write VastDB)
UI	Web app at your <code>INGRESS_URL</code> (frontend + backend)
S3 buckets	Chunks + segments for uploads
VastDB	Indexed segments, embeddings, detections, reasoning text

Piece	What it is
This repo in Cursor	Agent skills that know how to call every important API — open this project in Cursor and describe what you want
Source code repo	Full VSS Blueprint (<code>vss-blueprint</code>) — optional reading; the live stack is already up

`cd /vss-blueprint`

Cloning the Blueprint is **not a requirement**. The live UI, APIs, and Cursor skills in this repo are enough for a strong demo — search, filters, re-ingest, upload, dashboards, mini-apps on top of the archive. Do not rebuild or redeploy DataEngine ingest functions for the challenge.

Open the UI, log in with your team user, and open **this skills repo** in **Cursor**. Skills are how you move fast against the live APIs.

Video corpus (already indexed)

Your team’s VSS archive is **pre-ingested** from the lab corpus below. Explore and **re-ingest** with the prompt/metadata you need (`ingest/reingest-videos` / `reingest-chunk`).

One pipeline, one VastDB index, several operational lenses. Only `scenario` + metadata change. Anchor demo line:

Show me every clip, from any camera in any site, where a person is close to a moving vehicle

That should pull I-24 highway traffic, Toronto live driving, neighborhood / SF street cams, and warehouse forklift / aisle views into one result set.

Search UI filters (upload metadata): `camera_id` , `capture_type` , `location` , plus **object class** from detections. Corpus tables below also list a Category column for demo grouping — that is not a filter field. For re-ingest, set analysis `scenario` (presets: `surveillance` , `traffic` , `live_driving` , `retail` , `warehouse` , `egocentric` , `sports` , `nhl` , `general`).

Use-case groups

1. Highway Multi-Cam Traffic — traffic

Source	Location	Category	camera_id	What's indexed
I-24 / 3D traffic	nashville	Traffic	i24_cam-1	~51 multi-cam highway clips (scene*_p*c*)

Multiple cameras per scene. Story: congestion, lane behavior, vehicle–vehicle interaction on a highway corridor. Try: *“truck changing lanes”, “dense traffic on the highway”, “vehicle braking hard.”*

2. Live Driving & Road Safety — live_driving / streets

Source	Location	Category	camera_id	What's indexed
PIE drives	toronto	Streets	pie_cam-3	6 long drive sets (set01 ... set06)

Forward-facing / on-road drives. Story: road-scene understanding and driving safety. Try: *“pedestrian near the curb”, “intersection with turning traffic”, “vehicle ahead braking.”*

3. Neighborhood Street Surveillance — surveillance / streets

Source	Location	Category	camera_id	What's indexed
Neighborhood cars	neighborhood	Streets	neighborhood_cam-1	2 day merges (2026-09-01, 2026-09-02)

Residential / neighborhood car activity. Story: street-level vehicle events over time. Try: *“car passing in front of houses”, “two vehicles close together”, “vehicle stopping at the curb.”*

4. SF Street Surveillance — surveillance / streets

Source	Location	Category	camera_id	What's indexed
SF streets	san_francisco	Streets	sf_streets_cam-1 ... sf_streets_cam-4	4 street cameras (sf-streets/1 ... 4) — ingesting soon

Fixed San Francisco street cams. Story: urban street activity across multiple viewpoints. Try: *“pedestrians crossing while cars wait”, “busy intersection”, “same moment from two SF cameras.”*

5. Warehouse Safety & Operations — warehouse

Source	Location	Category	camera_id	What's indexed
SDG warehouse RGB	warehouse3	Warehouse	sdg_warehouse_cam-2	~178 short ceiling / aisle clips

Industrial **safety & near-miss** search — *“forklift near a person in an aisle”, “pallet in a walkway”, “person in a restricted zone.”*

6. Indoor / Smart-Space Surveillance — surveillance / crowds

Source	Location	Category	camera_id	What's indexed
Smart spaces	indoor	Crowds	smartspace_cam-1	~102 indoor / facility camera clips

Occupancy, people flow, and facility ops. Try: *“person walking through an aisle”, “group of people near equipment”, “empty corridor.”*

Cross-group demos (the payoff)

Query	Hits across groups
<i>“person close to a moving vehicle”</i>	I-24 (1) + Toronto driving (2) + neighborhood (3) + SF streets (4) + warehouse (5)
<i>“dense traffic / many vehicles”</i>	I-24 multi-cam (1) + PIE sets (2) + neighborhood (3) + SF streets (4)
<i>“person in a warehouse aisle”</i>	SDG RGB (5) + smartspaces (6)
<i>“same city, different camera”</i>	SF: sf_streets_cam-1 ... sf_streets_cam-4

Recommended demo packs

Don't try to demo everything at once. Start with a short pack, prove search → filters → UI, then expand.

Pack	Sources	Story
A — Highway Traffic	I-24 multi-cam	Corridor traffic
B — Live Driving	PIE sets 01–06	On-road / road-safety
C — Warehouse Safety	SDG RGB	Near-miss + aisle ops
D — Neighborhood Streets	Neighborhood day 01–02	Day-scale street vehicles
E — SF Streets	SF cams 1–4	Multi-cam urban streets (ingesting soon)
F — Indoor Smart Spaces	Smart spaces cams	Facility / indoor crowds

Start with **Pack A (Highway Traffic)** or **Pack C (Warehouse Safety)** — high-signal — before leaning on the long PIE / Neighborhood day merges. Add **Pack E (SF Streets)** once those cameras are indexed.

Example queries

- “Truck changing lanes on the highway” → i24_cam-1
- “Pedestrian near the road while driving” → pie_cam-3
- “Car passing houses on a residential street” → neighborhood_cam-1
- “Busy San Francisco intersection” → sf_streets_cam-1 ... sf_streets_cam-4
- “Forklift approaching a person in a warehouse aisle” → sdg_warehouse_cam-2
- “Person walking through an indoor corridor” → smartspace_cam-1
- “Person close to a moving vehicle” → cross-pack (A + B + C + D + E)

Suggested playbook

Write full prompts in Cursor: name what you want, point at your team config, and say what to build. A solid loop is **re-ingest → confirm indexing → ship a thin app on search / videos / dashboard**.

Example prompts (adapt to your use case)

1. Re-ingest and verify indexing

Re-ingest an indexed video from Pack A (Highway Traffic) or Pack C (Warehouse Safety) in my team environment (credentials in `/config/<my-team>.config`). Let me pick the target (e.g. an I-24 scene clip or an SDG warehouse clip), prompt/metadata behavior, and chunk count. Wait until re-ingest finishes, then run basic sanity checks on counts and pipeline health.
(Cursor will use skills: `ingest/reingest-videos`, `retrieval/dashboard`, `retrieval/login`)

2. Cross-camera “person near vehicle” board

Build a small webpage that answers: “*person close to a moving vehicle*” across my archive. Group hits by `location` / `camera_id` (Nashville I-24, Toronto driving, neighborhood streets, SF streets when indexed, warehouse). For each hit, show a triptych: one clip **before**, the **event**, and one clip **after** (~5 seconds each). Deploy it **on Kubernetes** (not local) to my team namespace without Docker build/push — Ingress path `/app` on my team host.
(Cursor will use skills: `retrieval/login`, `retrieval/search`, `retrieval/videos`, `retrieval/list-metadata`, `deployment/deploy-app-no-registry`)

3. Warehouse safety ops board

Build a standalone page for Pack C: histogram of top detected objects + a list of near-miss style hits (forklift near a person, tight aisle, person in a walkway). For each object, show a small bounding-box crop from a real segment. Filter by `camera_id` / `location`. Deploy on K8s at `/app` (or save under `tools/` while iterating), then report the top findings.
(Cursor will use skills: `retrieval/login`, `retrieval/dashboard`, `retrieval/search`, `retrieval/videos`, `deployment/deploy-app-no-registry`)

Rules of the road

- Stay in **your** team credentials, buckets, and UI. Don't poke other teams' namespaces.
- Prefer **Cursor + skills** over hand-copying curl forever — but reading a skill once to understand the API is encouraged.
- Don't burn the whole hackathon redeploying infrastructure; the stack is already up.

- If search returns nothing: check login, check dashboard/pipeline alignment, then check that metadata filter values actually exist (`retrieval/list-metadata`).
- Have fun — the win is a crisp story: *problem* → *video archive (one of the packs)* → *search/filters across cameras* → *insight or action*.